

Class 08: Breast Cancer

Kimberly Navarro (A17485724)

Table of contents

Background	1
Data Import	1
Principal Component Analysis (PCA)	5
Interpreting PCA results	6
Variance explained	8
Communicating PCA results	10
4. Hierarchical clustering	11
Selecting number of clusters	12
Using different methods	12
Combining methods	13
Sensitivity/ Specificity	14
Prediction	15

Background

In today's lecture we will be employing all the R techniques for data analysis that we have learned thus far - including the machine learning methods of clustering and PCA - to analyze real breast cancer biopsy data.

Data Import

The data is in CSV format:

```
wisc.df <- read.csv("WisconsinCancer (3).csv", row.names=1)
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001		0.14710
842517	0.08474	0.07864	0.0869		0.07017
84300903	0.10960	0.15990	0.1974		0.12790
84348301	0.14250	0.28390	0.2414		0.10520
84358402	0.10030	0.13280	0.1980		0.10430
843786	0.12780	0.17000	0.1578		0.08089
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
84300903	0.2069		0.05999	0.7456	0.7869
84348301	0.2597		0.09744	0.4956	1.1560
84358402	0.1809		0.05883	0.7572	0.7813
843786	0.2087		0.07613	0.3345	0.8902
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867
84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003		0.006193	25.38	17.33
842517	0.01389		0.003532	24.99	23.41
84300903	0.02250		0.004571	23.57	25.53
84348301	0.05963		0.009208	14.91	26.50
84358402	0.01756		0.005115	22.54	16.67
843786	0.02165		0.005082	15.47	23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.60	2019.0	0.1622		0.6656
842517	158.80	1956.0	0.1238		0.1866
84300903	152.50	1709.0	0.1444		0.4245
84348301	98.87	567.7	0.2098		0.8663
84358402	152.20	1575.0	0.1374		0.2050
843786	103.40	741.6	0.1791		0.5249
	concavity_worst	concave.points_worst	symmetry_worst		

842302	0.7119	0.2654	0.4601
842517	0.2416	0.1860	0.2750
84300903	0.4504	0.2430	0.3613
84348301	0.6869	0.2575	0.6638
84358402	0.4000	0.1625	0.2364
843786	0.5355	0.1741	0.3985
	fractal_dimension_worst		
842302		0.11890	
842517		0.08902	
84300903		0.08758	
84348301		0.17300	
84358402		0.07678	
843786		0.12440	

peak at the data

```
head(wisc.df, 3)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.8	1001
842517	M	20.57	17.77	132.9	1326
84300903	M	19.69	21.25	130.0	1203
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001		0.14710
842517	0.08474	0.07864	0.0869		0.07017
84300903	0.10960	0.15990	0.1974		0.12790
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
84300903	0.2069		0.05999	0.7456	0.7869
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003		0.006193	25.38	17.33
842517	0.01389		0.003532	24.99	23.41
84300903	0.02250		0.004571	23.57	25.53
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.6	2019	0.1622		0.6656
842517	158.8	1956	0.1238		0.1866
84300903	152.5	1709	0.1444		0.4245

```

concavity_worst concave.points_worst symmetry_worst
842302          0.7119           0.2654        0.4601
842517          0.2416           0.1860        0.2750
84300903        0.4504           0.2430        0.3613
fractal_dimension_worst
842302            0.11890
842517            0.08902
84300903        0.08758

```

Q1. How many observations are in this dataset?

There are 569 observations of 31 variables

```
nrow(wisc.df)
```

```
[1] 569
```

Q2. How many of the observations have a malignant diagnosis?

212 malignant diagnosis

```
table(wisc.df$diagnosis)
```

	B	M
357	212	

Q3. How many variables/features in the data are suffixed with _mean?

There are 10 variables/feaures suffixed with _mean

```
length(grep("_mean", colnames(wisc.df)) )
```

```
[1] 10
```

We need to remove the **diagnosis** column before we do any further analysis of this dataset - we don't want to pass this to PCA etc. We will save it as a separate wee vector that we can use later to compare our findings to those of experts.

```
wisc.data <- wisc.df[ , -1]
diagnosis <- wisc.df$diagnosis
```

Principal Component Analysis (PCA)

The main function in base R is called `prcomp()` we will use the optional argument `scale=TRUE` here as the data columns/features/dimensions are on very different scales in the original set.

```
wisc.pr <- prcomp(wisc.data, scale=T)
```

Inspection of summary

```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

```
summary(wisc.pr)$importance[2,1]
```

[1] 0.44272

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

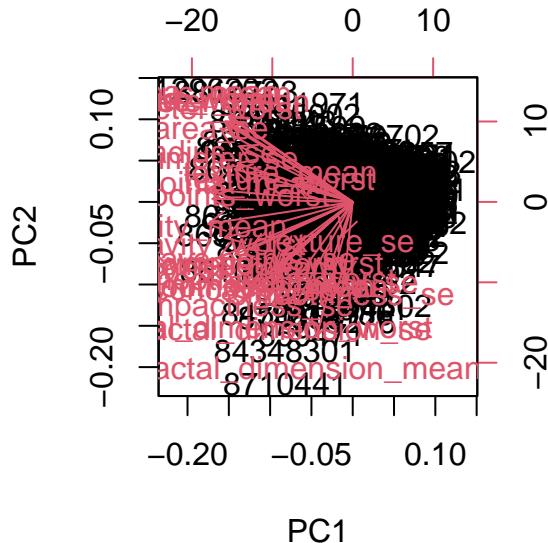
In order to describe at least 70% of the original variance in the data, there must be 3 PCs.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

In order to describe at least 90% of the original variance in the data, there must be 7 PCs.

Interpreting PCA results

```
biplot(wisc.pr)
```



Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

The plot is extremely difficult to understand as everything is overlapping. Row names are being used as plotting characters.

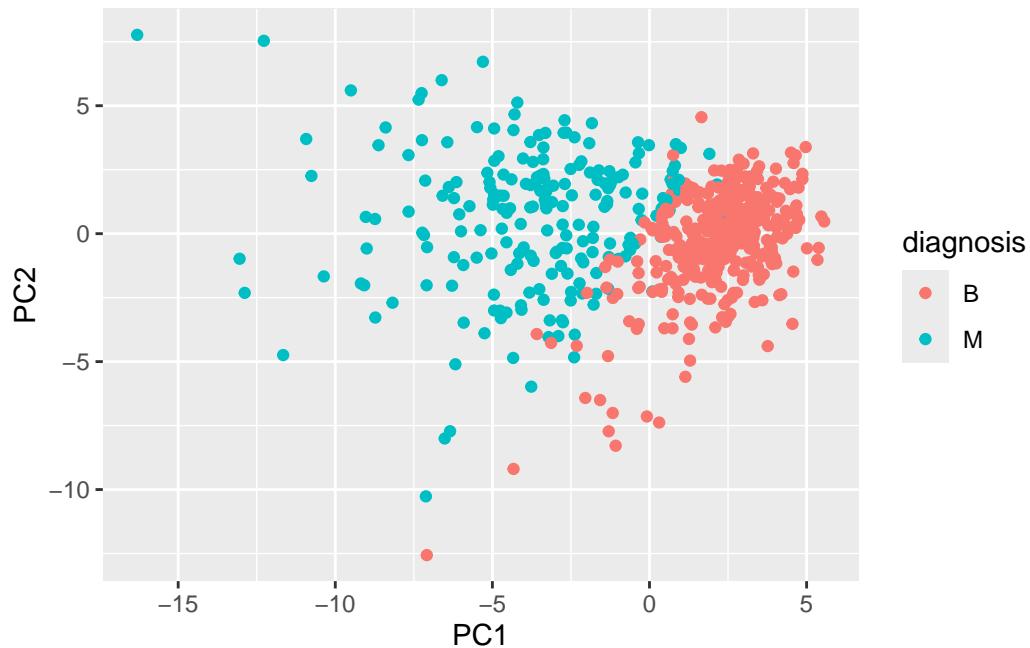
Lets generate a more standard scatter plot of each observation along components 1&2:

```

library(ggplot2)

ggplot(wisc.pr$x) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()

```



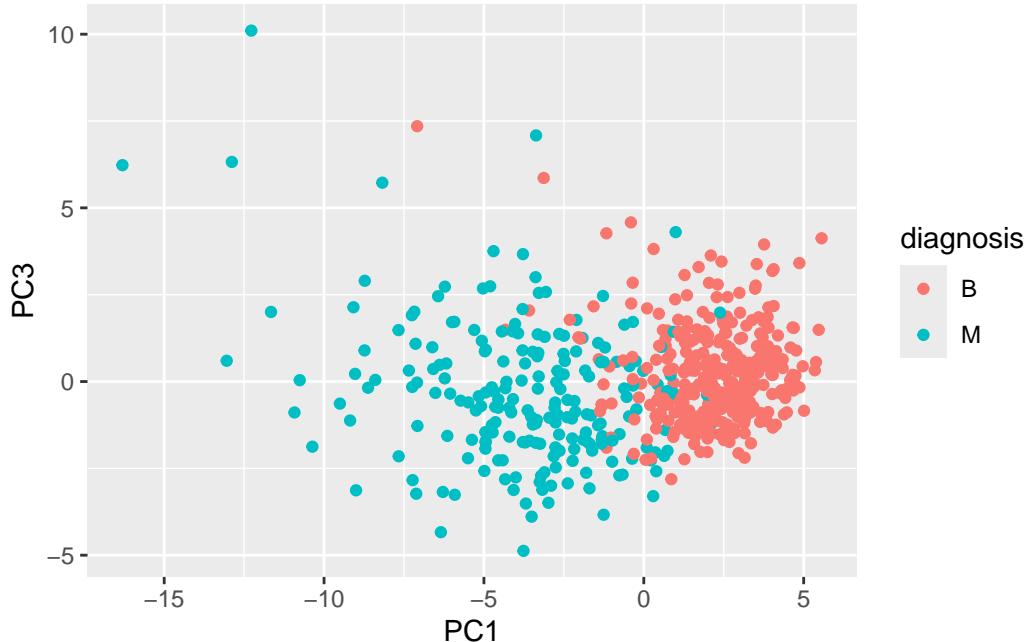
Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

The plots indicate that P1 is capturing a separation of malignancy (blue) from benign samples (red). As shown in P1 & P3 below:

```

ggplot(wisc.pr$x) +
  aes(PC1, PC3, col=diagnosis) +
  geom_point()

```



Variance explained

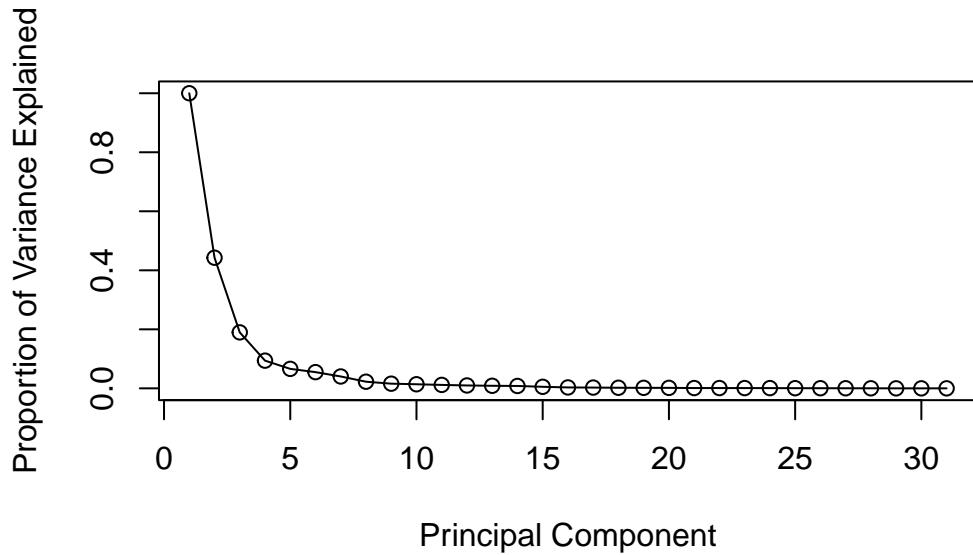
A scree plot shows how much variance each PC captures. The “elbow”- point where adding more PCs gives diminishing returns can help us decide how many PCs should we consider for further analysis.

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

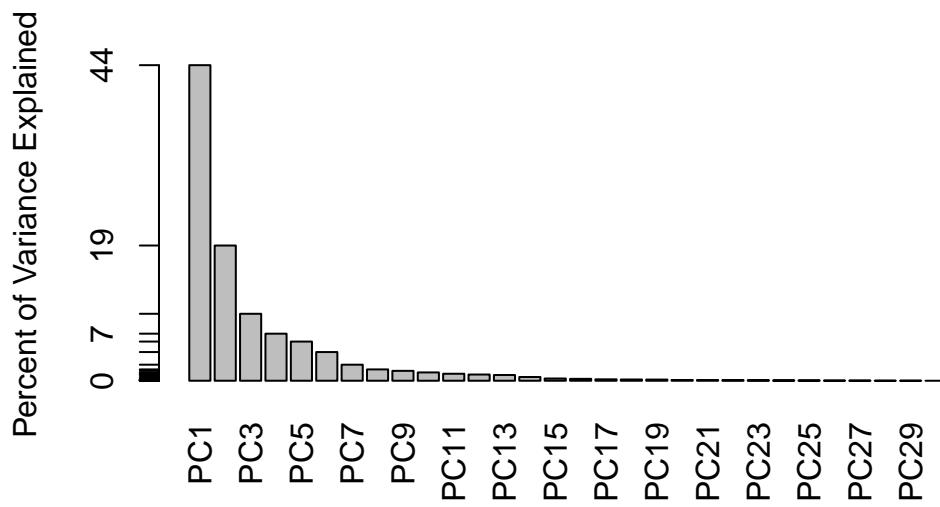
```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

```
pve <- (wisc.pr$sdev^2)/(ncol(wisc.pr$x))

plot(c(1,pve), xlab = "Principal Component",
      ylab = "Proportion of Variance Explained",
      ylim = c(0, 1), type = "o")
```



```
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```

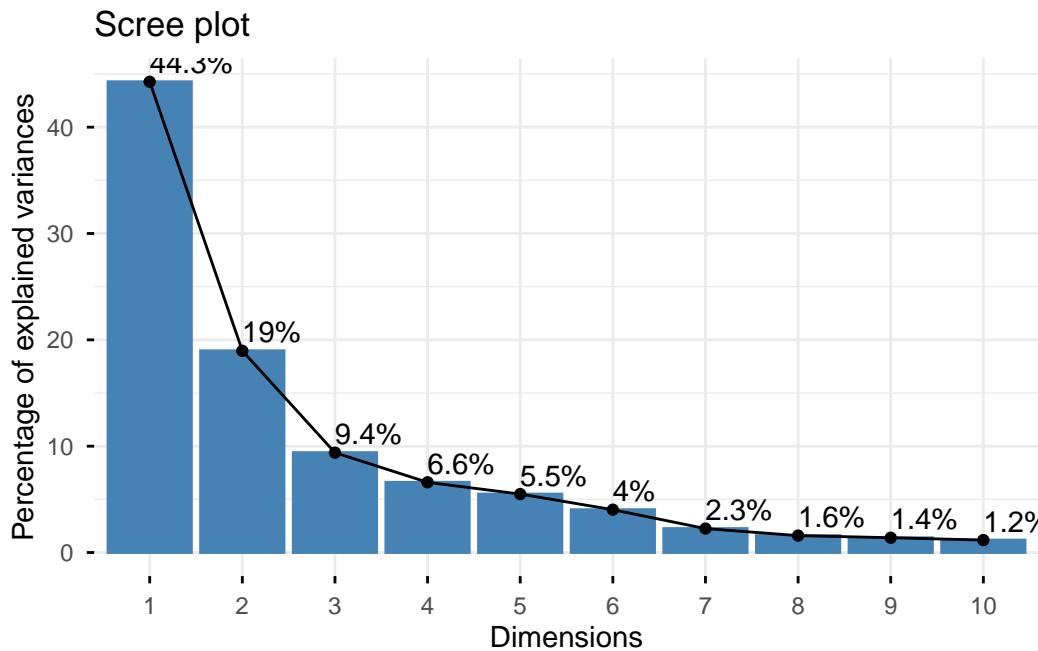


```
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
fviz_eig(wisc.pr, addlabels = TRUE)
```

Warning in geom_bar(stat = "identity", fill = barfill, color = barcolor, :
Ignoring empty aesthetic: `width`.



Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. wisc.pr\$rotation[,1]) for the feature concave.points_mean?

```
wisc.pr$rotation["concave.points_mean", 1]
```

[1] -0.2608538

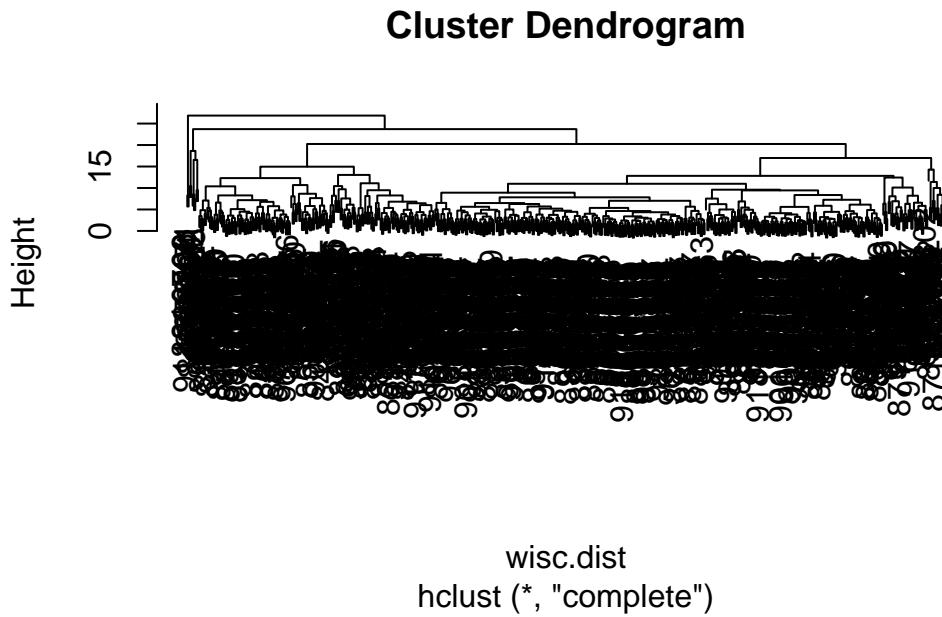
4. Hierarchical clustering

The goal of this section is to do hierarchical clustering of the original data to see if there is any obvious grouping into malignant and benign clusters.

In short, these results are not good!

First, we will scale our `wisc.data` then calculate a distance matrix, then pass to `hclust()`:

```
wisc.dist <- dist( scale(wisc.data) )
wisc.hclust <- hclust(wisc.dist)
plot(wisc.hclust)
```

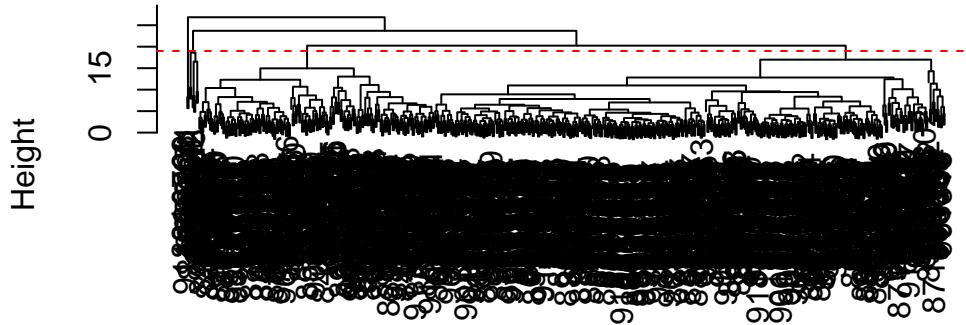


Q10. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

At approximately $h=19$, is where clustering model has 4 clusters.

```
plot(wisc.hclust)
abline(h = 19, col="red", lty=2)
```

Cluster Dendrogram



```
wisc.dist  
hclust (*, "complete")
```

Selecting number of clusters

```
wisc.hclust.clusters <- cutree(wisc.hclust, k=2)  
table(wisc.hclust.clusters)
```

```
wisc.hclust.clusters  
1 2  
567 2
```

Using different methods

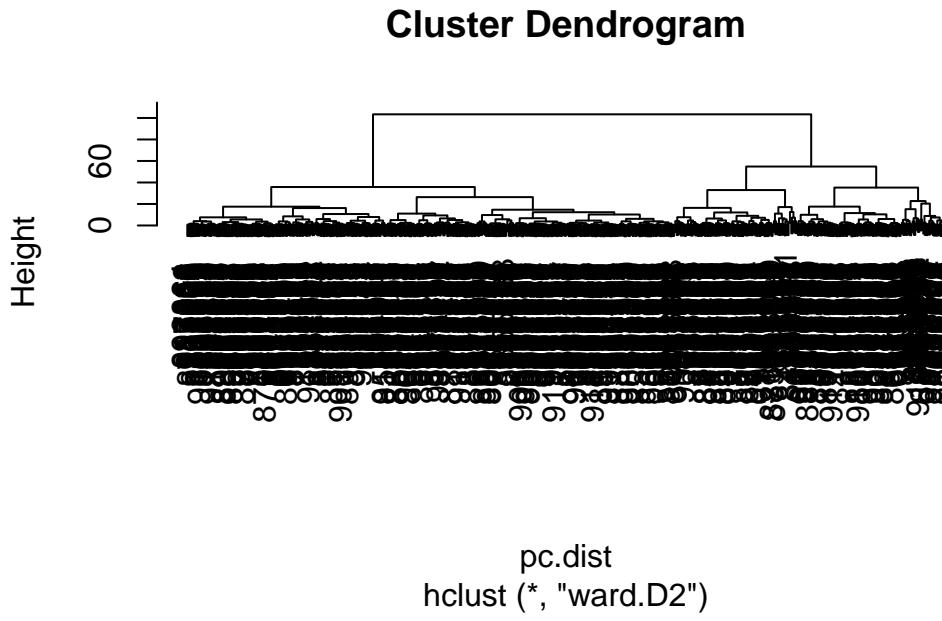
Q12. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning

Ward.D2 gives the most structured and separated results, making them easy to interpret.

Combining methods

The idea here is that I can take my new variables (i.e. the scores on the PCs `wisc.pr$x`) that are better descriptors of the data-set than the original features (i.e. the 30 columns in `wisc.data`) and use these as a basis for clustering.

```
pc.dist <- dist(wisc.pr$x[, 1:3])
wisc.pr.hclust <- hclust(pc.dist, method="ward.D2")
plot(wisc.pr.hclust)
```



We can “cut” this tree to yield our clusters:

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
 1   2
203 366
```

```
table(diagnosis)
```

```
diagnosis
  B   M
357 212
```

Q13. How well does the newly created hclust model with two clusters separate out the two “M” and “B” diagnoses?

The newly created hclust model with 2 clusters does a decent job at separating the “M” and “B” diagnosis.

I can now run `table()` with both my clustering `grps` and the expert `diagnosis`

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses?

The hierachial clustering models created in previous sections do not perform a good job at separating diagnosis. The PCA variables allow for a better separation as shown below:

```
table(grps, diagnosis)
```

```
diagnosis
grps   B   M
  1   24 179
  2 333  33
```

Sensitivity/ Specificity

Our cluster “1” has 179 “M” diagnosis Our cluster “2” has 333 “B” diagnosis

179 TP 24 TP 333 TN 33 FN

Sensitivity: $TP/(TP+FN)$

```
179/(179+33)
```

```
[1] 0.8443396
```

Specificity: $TN/(TN+FP)$

```
333/(333+24)
```

```
[1] 0.9327731
```

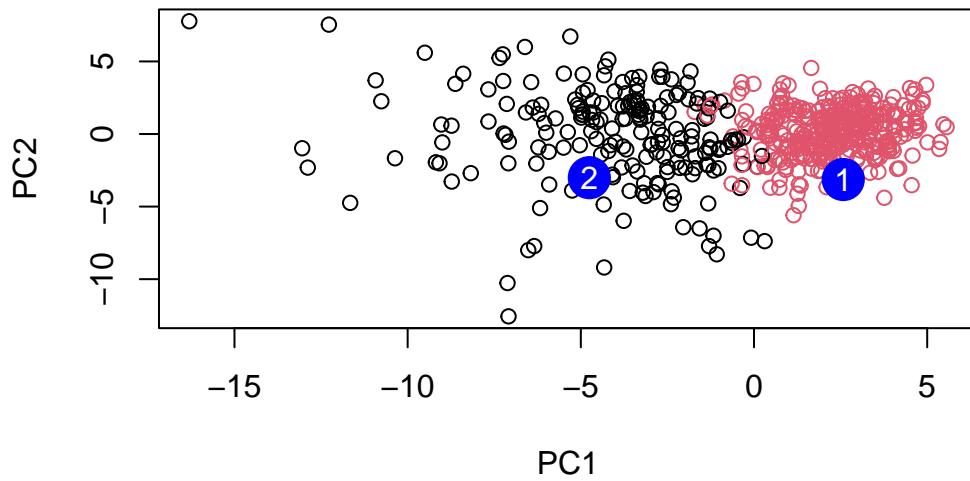
Prediction

We can use our PCA model for prediction of new un-seen causes:

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

```
PC1      PC2      PC3      PC4      PC5      PC6      PC7
[1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
[2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
      PC8      PC9      PC10     PC11     PC12     PC13     PC14
[1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
[2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
      PC15     PC16     PC17     PC18     PC19     PC20
[1,]  0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
[2,]  0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
      PC21     PC22     PC23     PC24     PC25     PC26
[1,]  0.1228233 0.09358453 0.08347651  0.1223396  0.02124121 0.078884581
[2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
      PC27     PC28     PC29     PC30
[1,]  0.220199544 -0.02946023 -0.015620933 0.005269029
[2,] -0.001134152 0.09638361  0.002795349 -0.019015820
```

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q.16 Which of these new patients should we prioritize for follow up based on your results?

Patient 1 should be prioritized for follow up as there seems to be more clustering with the malignant cases (cancerous) as shown in red.